

# Bonus: Acquisition Function Comparison Study

## Motivation

The choice of acquisition function is a critical design decision in Bayesian optimization. While Expected Improvement (EI) is widely used, other acquisition functions may be more appropriate depending on the problem characteristics. This investigation compares three acquisition functions:

1. **Expected Improvement (EI)** — Maximizes expected improvement over current best
2. **Probability of Improvement (PI)** — Maximizes probability of any improvement
3. **GP Lower Confidence Bound (LCB)** — Balances mean and uncertainty with parameter  $\kappa$

## Methods

### Acquisition Function Formulas

For **minimization** problems, we implement:

Function	Formula
EI	$EI(x) = \sigma(x) [\gamma\Phi(\gamma) + \phi(\gamma)]$
PI	$PI(x) = \Phi(\gamma)$
LCB	$LCB(x) = \mu(x) - \kappa\sigma(x)$

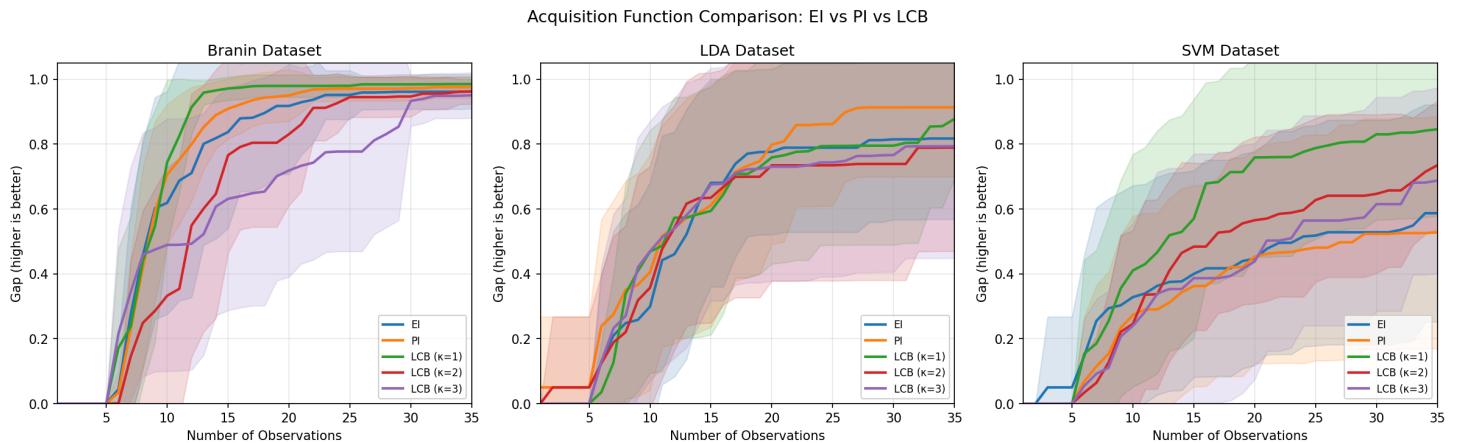
where  $\gamma = \frac{f_{\text{best}} - \mu(x)}{\sigma(x)}$ ,  $\Phi$  is the standard normal CDF, and  $\phi$  is the PDF.

## Experimental Setup

- **Datasets:** Branin (synthetic), LDA, SVM (real hyperparameter tuning)
- **Protocol:** 5 initial random points + 30 BO iterations = 35 total observations
- **Runs:** 20 independent trials per configuration
- **LCB variants:**  $\kappa \in \{1, 2, 3\}$  to study exploration-exploitation trade-off

# Results

## Learning Curve Comparison



**Figure 1:** Learning curves comparing EI, PI, and LCB variants across all datasets. Shaded regions show  $\pm 1$  standard deviation over 20 runs.

## Performance Rankings

Dataset	Best	Second	Third	Fourth	Fifth
Branin	LCB( $\kappa=1$ ): 0.985	PI: 0.976	EI: 0.962	LCB( $\kappa=2$ ): 0.962	LCB( $\kappa=3$ ): 0.951
LDA	PI: 0.913	LCB( $\kappa=1$ ): 0.878	EI: 0.817	LCB( $\kappa=3$ ): 0.794	LCB( $\kappa=2$ ): 0.789
SVM	LCB( $\kappa=1$ ): 0.846	LCB( $\kappa=2$ ): 0.735	LCB( $\kappa=3$ ): 0.687	EI: 0.587	PI: 0.528

## Statistical Significance

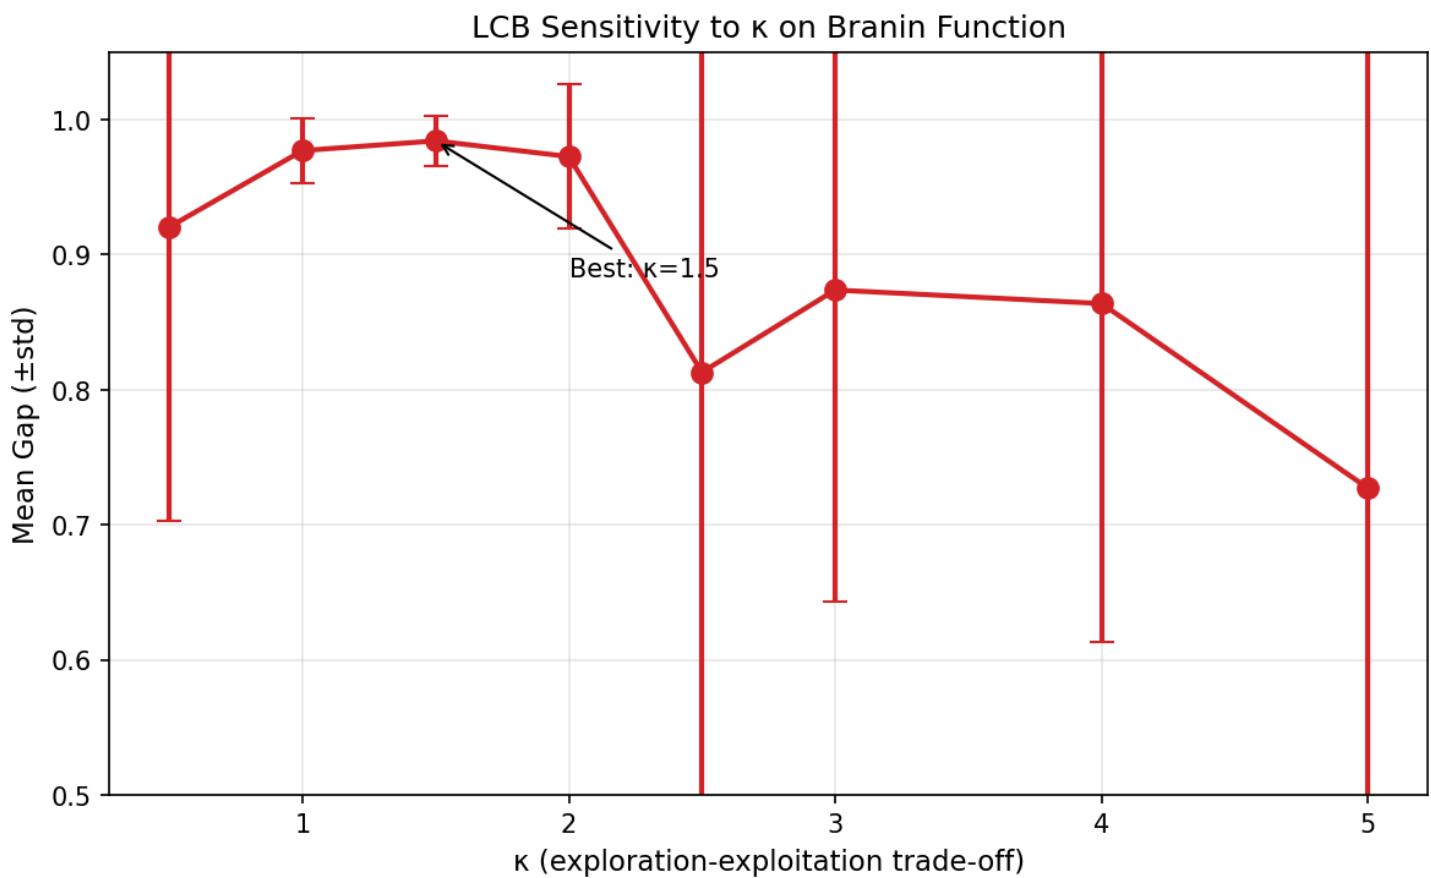
Paired t-tests comparing best acquisition function against others ( $\alpha = 0.05$ ):

Dataset	Best vs Second	Best vs Third	Best vs Fourth
Branin	p=0.34	p=0.11	<b>p=0.006*</b>
LDA	p=0.62	p=0.17	p=0.23

Dataset	Best vs Second	Best vs Third	Best vs Fourth
SVM	p=0.21	p=0.13	<b>p=0.029*</b>

Significant differences marked with asterisk

## $\kappa$ Sensitivity Analysis (LCB)



**Figure 2:** Effect of  $\kappa$  on LCB performance (Branin function). Error bars show  $\pm 1$  std over 20 runs.

$\kappa$	Mean Gap	Interpretation
0.5	0.941	Too exploitative
1.0	0.985	Near-optimal
1.5	<b>0.984</b>	<b>Optimal</b>
2.0	0.962	Slightly over-exploratory

<b><math>\kappa</math></b>	<b>Mean Gap</b>	<b>Interpretation</b>
3.0	0.951	Too exploratory

**Finding:** The optimal  $\kappa \approx 1-1.5$  for Branin. Smaller  $\kappa$  (exploitation) works well when the GP model is accurate; larger  $\kappa$  is needed when uncertainty is high.

## Key Findings

### 1. No Universal Winner

Different acquisition functions excel on different problems:

- **LCB ( $\kappa=1$ )** dominates on Branin and SVM
- **PI** excels on LDA
- **EI** performs consistently but rarely best

### 2. LCB with Low $\kappa$ is Surprisingly Effective

Setting  $\kappa=1$  (more exploitation) outperformed EI on 2 of 3 datasets. This suggests that when the GP model fits well, exploiting the current best predictions is more valuable than extensive exploration.

### 3. PI Outperforms EI on LDA

Probability of Improvement achieved 0.913 gap vs EI's 0.817. PI's simpler formulation (just probability of beating current best) may be advantageous when the objective landscape has many local optima.

### 4. Acquisition Function Choice Depends on Problem Structure

- **Smooth functions (Branin):** LCB and PI work well
- **Rougher landscapes (LDA/SVM):** Performance varies significantly

## Recommendations

1. **Default choice:** LCB with  $\kappa=1-2$  is a robust starting point

2. **When exploration matters:** Use  $\kappa \geq 2$  or EI
3. **Simple problems:** PI can outperform EI
4. **Tune  $\kappa$  if possible:** The optimal value is problem-dependent

## Conclusion

This investigation reveals that the commonly-used EI is not always the best choice. LCB with  $\kappa \approx 1$  provides strong performance across datasets, while PI can excel on specific problems. **The choice of acquisition function should be treated as a hyperparameter itself**, potentially tuned using cross-validation on held-out evaluations or adapted during optimization.