

# Data Visualization

This section presents our initial exploration of the objective functions used in this Bayesian optimization investigation: the synthetic Branin benchmark and the real-world hyperparameter tuning datasets for SVM and LDA models.

## Bullet Point Reference


This report addresses each bullet point from the [instructions](#):

Bullet	Instruction Summary	Report Section
1	Make a $1000 \times 1000$ heatmap of Branin function	<a href="#">Section 1</a>
2	Describe behavior and stationarity	<a href="#">Section 2</a>
3	Find transformation for stationarity	<a href="#">Section 3</a>
4	KDE of LDA and SVM distributions, interpret	<a href="#">Section 4</a>
5	Find transformation for better-behaved distributions	<a href="#">Section 5</a>

## 1. Branin Function Heatmap

**Bullet 1:** "Make a heatmap of the value of the Branin function over the domain  $X = [-5, 10] \times [0, 15]$  using a dense grid of values, with 1000 values per dimension, forming a  $1000 \times 1000$  image."

We evaluated the Branin function over the domain  $\mathcal{X} = [-5, 10] \times [0, 15]$  using a dense  $1000 \times 1000$  grid (1,000,000 evaluation points).



**Figure 1:** Heatmap of the Branin function. Red stars mark the three global minima at  $(-\pi, 12.275)$ ,  $(\pi, 2.275)$ , and  $(9.42478, 2.475)$ , each with value  $f^* \approx 0.398$ .

The Branin function is defined as:

$$f(x_1, x_2) = a (x_2 - bx_1^2 + cx_1 - r)^2 + s(1 - t) \cos(x_1) + s$$

with standard parameters  $a = 1$ ,  $b = \frac{5.1}{4\pi^2}$ ,  $c = \frac{5}{\pi}$ ,  $r = 6$ ,  $s = 10$ ,  $t = \frac{1}{8\pi}$ .

### Key observations:

- Function values range from approximately 0.4 (at minima) to over 300 (at corners)
- Three distinct global minima create multiple attraction basins
- The parabolic valley structure is modulated by the cosine term

## 2. Stationarity Analysis

**Bullet 2:** "Describe the behavior of the function. Does it appear stationary? (That is, does the behavior of the function appear to be relatively constant throughout the domain?)"

**Does it appear stationary?** No. **The Branin function is non-stationary.** This is evident from several observations:

1. **Varying Magnitude:** The function values span a dramatic range from 0.4 to 308, indicating non-constant behavior across the domain.
2. **Asymmetric Structure:** The three global minima are not uniformly distributed—two lie near  $x_2 \approx 2.5$  while one is at  $x_2 \approx 12.3$ , creating an asymmetric landscape.
3. **Quadratic Component:** The term  $a(x_2 - bx_1^2 + cx_1 - r)^2$  creates a parabolic valley whose curvature varies with position.
4. **Periodic Modulation:** The cosine term  $s(1 - t) \cos(x_1)$  adds oscillation in the  $x_1$  direction only, creating wave-like patterns that interact with the quadratic structure.
5. **Edge Effects:** Function values are significantly higher near domain boundaries, especially at the corners.

### Implications for Bayesian Optimization:


- A stationary GP prior may struggle to model this varying behavior
- The optimizer may require more samples in high-variance regions
- Non-stationary or adaptive kernels may improve surrogate model quality

## 3. Transformation for Improved Stationarity

**Bullet 3:** "Can you find a transformation of the data that makes it more stationary?"

**Yes.** We apply a log transformation to compress the dynamic range and improve stationarity:

$$g(x_1, x_2) = \log(f(x_1, x_2) + 1)$$



**Figure 2:** Comparison of original (left) and log-transformed (right) Branin function.

Metric	Original	Log-Transformed
Range	[0.40, 308.13]	[0.34, 5.73]
Max/Min Ratio	~750x	~17x
Std. Deviation	51.35	1.12

### Why does this help stationarity?


- Compresses the dynamic range by a factor of ~44x
- Reduces extreme values at domain boundaries
- Preserves the locations of minima and overall structure
- Provides approximate variance stabilization
- Makes the function appear more uniform across the domain

## 4. Kernel Density Estimates for LDA and SVM Benchmarks

**Bullet 4:** "Make a kernel density estimate of the distribution of the values for the LDA and SVM benchmarks. Interpret the distributions."

We analyze the distribution of objective values for the hyperparameter tuning benchmarks:

- **LDA:** 288 hyperparameter configurations
- **SVM:** 1,400 hyperparameter configurations



**Figure 3:** Kernel density estimates of the objective value distributions. Dashed red lines indicate the minimum (best) performance.

## LDA Benchmark

Statistic	Value
Samples	288
Range	[1266.17, 5258.11]
Mean	1820.67
Std. Dev.	722.33
Skewness	2.37 (right-skewed)

**Interpretation:** The LDA objective exhibits a right-skewed distribution with a long tail toward poor performance. Most configurations cluster around moderate values (1300–2000), while some produce very poor results (>3000). Optimal configurations are relatively rare.

## SVM Benchmark

Statistic	Value
Samples	1,400

Statistic	Value
Range	[0.2411, 0.5000]
Mean	0.3136
Std. Dev.	0.0693
Skewness	1.30 (right-skewed)

**Interpretation:** The SVM error rate clusters around 0.27–0.35, with outliers at 0.5 (random chance). The pronounced mode suggests many configurations achieve similar moderate performance. The minimum (~0.24) represents near-optimal classification accuracy.

**Common Pattern:** Both distributions are right-skewed, reflecting a common property of hyperparameter landscapes—there are many ways to configure a model poorly but relatively few optimal configurations.


## 5. Transformation for Better-Behaved Distributions

**Bullet 5:** "Again, can you find a transformation that makes the performance better behaved?"

**Yes.** We apply a log transformation to both benchmark distributions to improve their statistical properties:

$$y' = \log(y)$$

Comparison: Original vs Log-Transformed Distributions

**Figure 4:** Comparison of original (left) and log-transformed (right) distributions for both benchmarks.

## Skewness Reduction

Benchmark	Original Skewness	Log-Transformed	Reduction
LDA	2.368	1.353	<b>42.9%</b>
SVM	1.302	1.078	<b>17.2%</b>

**Why does log transformation make distributions "better behaved"?**

- Improved Symmetry:** Both distributions become more symmetric (closer to Gaussian), which is beneficial for GP likelihood assumptions.
- Variance Stabilization:** The variance becomes more uniform across the range of values.
- Outlier Compression:** Extreme high values are compressed, reducing their influence on the GP fit.
- Normal Approximation:** Log-transformed data better approximates a Gaussian distribution.

**Recommendation:** Use log-transformed objective values when fitting Gaussian processes for Bayesian optimization on these benchmarks.

## Summary

Bullet	Question	Answer
1	Heatmap created?	Yes - 1000×1000 grid with 3 marked global minima
2	Is Branin stationary?	<b>No</b> - varies from 0.4 to 308, asymmetric structure
3	Transformation for stationarity?	<b>log(f+1)</b> reduces dynamic range 44×
4	KDE interpretation?	Both LDA and SVM are <b>right-skewed</b>
5	Transformation for better behavior?	<b>log(y)</b> reduces skewness by 17–43%